

YI DING

BHEE 336, 465 Northwestern Ave ◊ West Lafayette, IN 47907, USA

Email: yiding@purdue.edu ◊ Website: www.y-ding.github.io

RESEARCH INTERESTS

Sustainable Computing, Machine Learning for Systems, Datacenter Computing, Causal Inference

PROFESSIONAL EXPERIENCE

Purdue University Assistant Professor in Elmore Family School of Electrical and Computer Engineering PI, STYLE (SusTainable computing sYstems and LEarning) Lab	West Lafayette, IN, USA 8/2023 – Present
Massachusetts Institute of Technology Postdoctoral Associate & NSF Computing Innovation Fellow. Mentor: Michael Carbin	Cambridge, MA, USA 1/2021 – 8/2023
Meta Visiting Researcher	Cambridge, MA, USA 10/2021–12/2022
Google Research Intern	Sunnyvale, CA, USA 6/2019–9/2019

EDUCATION

University of Chicago Ph.D. & MS in Computer Science. Advisor: Henry Hoffmann	Chicago, IL, USA 8/2015 – 12/2020
Nanyang Technological University Doctoral Student in Computer Science. Passed Qualification Exam.	Singapore 7/2013 – 7/2015
Beijing Jiaotong University B.E. in Electronic Science and Technology. Graduated with Highest Honor.	Beijing, China 9/2008 – 6/2012

SELECTED AWARDS AND HONORS

CRA/CCC/NSF Computing Innovation Fellowship	2020-2023
Meta Research Award	2021
EECS Rising Stars at UC Berkeley	2020

PUBLICATIONS

- Uncertainty-Aware Decarbonization for Datacenters**
Amy Li, Sihang Liu, Yi Ding
Workshop on Sustainable Computer Systems (HotCarbon), 2024
- Towards Sustainable Large Language Model Serving**
Sophia Nguyen*, Beihao Zhou*, Yi Ding, Sihang Liu (*Equal Contributions)
Workshop on Sustainable Computer Systems (HotCarbon), 2024
- Turaco: Complexity-Guided Data Sampling for Training Neural Surrogates of Programs**
Alex Renda, Yi Ding, Michael Carbin
ACM SIGPLAN conference on Systems, Programming, Languages, and Applications (OOPSLA), 2023
- CAFQA: A Classical Simulation Bootstrap for Variational Quantum Algorithms**
Gokul Ravi, Pranav Gokhale, Yi Ding, William M. Kirby, Kaitlin N. Smith, Jonathan M. Baker, Peter J. Love, Henry Hoffmann, Kenneth R. Brown, Frederic T. Chong
ACM International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS), 2023
- Minimax Designs for Causal Effects in Temporal Experiments with Treatment Habituation**
Guillaume Basse, Yi Ding, Panos Toulis
Biometrika, (Top theoretical statistics journal), 2023

NURD: Negative-Unlabeled Learning for Online Datacenter Straggler Prediction

Yi Ding, Avinash Rao, Hyebin Song, Rebecca Willett, Henry Hoffmann

Conference on Machine Learning and Systems (MLSys), 2022**Generalizable and Interpretable Learning for Configuration Extrapolation**

Yi Ding, Ahsan Pervaiz, Michael Carbin, Henry Hoffmann.

ACM Joint European Software Engineering Conference and Symposium on the Foundations of Software Engineering (ESEC/FSE), 2021**Programming with Neural Surrogates of Programs**

Alex Renda, Yi Ding, Michael Carbin

ACM SIGPLAN International Symp. on New Ideas, New Paradigms, and Reflections on Programming and Software (Onward!), 2021**Neighborhood Street Activity and Greenspace Usage Uniquely Contribute to Predicting Crime**Kathryn Schertz, James Saxon, Carlos Cardenas-Iniguez, Luís Bettencourt, Yi Ding, Henry Hoffmann, Marc G Berman*npj Urban Sustainability, Nature Research Journal*, 2021**Dynamical Systems Theory for Causal Inference with Application to Synthetic Control Methods**

Yi Ding, Panos Toulis

International Conference on Artificial Intelligence and Statistics (AISTATS), 2020**A Polynomial-time Algorithm for Learning Nonparametric Causal Graphs**

Ming Gao, Yi Ding, Bryon Aragam

Advances in Neural Information Processing Systems (NeurIPS), 2020**Generative and Multi-phase Learning for Computer Systems Optimization**

Yi Ding, Nikita Mishra, Henry Hoffmann

International Symposium on Computer Architecture (ISCA), 2019**Multiresolution Kernel Approximation for Gaussian Process Regression**

Yi Ding, Risi Kondor, Jonathan Eskreis-Winkler

Advances in Neural Information Processing Systems (NeurIPS), 2017 (**Spotlight**)**Large Scale Kernel Methods for Online AUC Maximization**

Yi Ding, Chenghao Liu, Peilin Zhao, Steven CH Hoi

IEEE International Conference on Data Mining (ICDM), 2017 (**Long Oral**)**An Adaptive Gradient Method for Online AUC Maximization**

Yi Ding, Peilin Zhao, Steven CH Hoi, Yew-Soon Ong

AAAI Conference on Artificial Intelligence (AAAI), 2015 (**Oral**)**Learning Relative Similarity by Stochastic Dual Coordinate Ascent**

Pengcheng Wu, Yi Ding, Peilin Zhao, Chunyan Miao, Steven CH Hoi

AAAI Conference on Artificial Intelligence (AAAI), 2014

WORKSHOP CONTRIBUTIONS

Causal and Interpretable Learning for Datacenter Latency Prediction

Yi Ding, Avinash Rao, Henry Hoffmann

Women in Machine Learning Workshop co-located with NeurIPS (WiML), 2020**A Polynomial-time Algorithm for Learning Nonparametric Causal Graphs**

Ming Gao, Yi Ding, Bryon Aragam

Women in Machine Learning Workshop co-located with NeurIPS (WiML), 2020**Minimax Crossover Designs**

Yi Ding, Guillaume Basse, Panos Toulis

NeurIPS Workshop on “Do the right thing”: machine learning and causal inference for improved decision making (CausalML), 2019

Minimax Crossover Designs for Digital Experimentation

Guillaume Basse, Yi Ding, Panos Toulis

Conference on Digital Experimentation at MIT (CODE@MIT), 2019**Generative and Multi-phase Learning for Computer Systems Optimization**

Yi Ding, Nikita Mishra, Henry Hoffmann

Women in Machine Learning Workshop co-located with NeurIPS (WiML), 2019**Nonparametric Causal Inference in Dynamical Systems with Synthetic Controls**

Yi Ding, Panos Toulis

Women in Machine Learning Workshop co-located with NeurIPS (WiML), 2018

RESEARCH ADVISING

PhD Students

Tianyao Shi, Purdue University

Fall 2024–

William Meng, University of Pennsylvania

Fall 2022–

Master Students

Ashutosh Sharma, UIUC

Spring 2024–

Hyunji Kim, MIT

2021–2022

Undergraduate Students

Sarah Deniz, Purdue University (DUIRI)

Fall 2024–

Gavin Fortwendel, Purdue University (DUIRI)

Fall 2024–

Yuqi Bai, University of Waterloo

Fall 2024–

Zihan Pan, University of Waterloo

Fall 2024–

Amy Li, University of Waterloo (One HotCarbon'24 paper published)

Spring 2024

Beihao Zhou, University of Waterloo (One HotCarbon'24 paper published)

Spring 2024

Sophia Nguyen, University of Waterloo (One HotCarbon'24 paper published)

Spring 2024

Avinash Rao, University of Chicago (One MLSys'22 paper published)

2019–2020

GRANTS

Title: Conference: DESC: Type III: A Holistic AI Computing Framework: Incorporating the Water and Biodiversity Dimensions of Sustainability

Funder: NSF

Duration: 2024–2025

People: Inez Hua (PI), Yi Ding (co-PI)

Awarded: \$9,9992 (My share: 50%)

Title: Computing Innovation Fellows 2020 Project

Funder: NSF

Duration: 2020–2023

People: Michael Carbin (PI), Yi Ding

Awarded: \$295,704

Title: Meta Research Award on Statistics for Improving Insights, Models, & Decisions

Funder: Meta

Duration: 2021–2022

People: Michael Carbin (PI), Yi Ding (co-PI)

Awarded: \$46,000

TEACHING

Instructor, Purdue University, West Lafayette, IN

Machine Learning in Cloud Computing (ECE 69500)

Fall 2024

Python for Data Science (ECE 20875)

Spring 2024

Python for Data Science (ECE 20875)

Fall 2023

Teaching Assistant, University of Chicago, Chicago, IL
 Machine Learning and Large Scale Data Analysis (CMSC 25025)
 Machine Learning (CMSC 25400)
 Machine Learning (MPCS 53111)
 Machine Learning for Public Policy (CAPP 30255)

Spring 2017
 Winter 2017
 Spring 2016
 Winter 2016

PROFESSIONAL SERVICE

Program Committee

IEEE International Symposium on High-Performance Computer Architecture (HPCA)	2025
USENIX Annual Technical Conference (ATC)	2024
Conference on Systems and Machine Learning (MLSys)	2024
ACM Student Research Competition at PACT	2023
SPLASH Onward!	2022
Conference on Systems and Machine Learning (MLSys)	2022
ACM Asia-Pacific Workshop on Systems	2022
Journal of Systems Research	2022

Technical Reviewing

Neural Information Processing Systems (NeurIPS)	2022
International Conference on Learning Representations (ICLR)	2022
International Conference on Machine Learning (ICML)	2022
Neural Information Processing Systems (NeurIPS)	2021
AAAI Conference on Artificial Intelligence (AAAI)	2021
AAAI Conference on Artificial Intelligence (AAAI)	2020
Neural Information Processing Systems (NeurIPS)	2019
International Conference on Machine Learning (ICML)	2019

INVITED TALKS

A Holistic View on Machine Learning for Systems

University of Waterloo, Department of Computer Science	Jun. 2023
Microsoft Research	Apr. 2023
Texas A&M University, Department of Computer Science & Engineering	Apr. 2023
University of Southern California, Department of Electrical & Computer Engineering	Apr. 2023
University of Illinois, Department of Computer Science	Mar. 2023
Cornell Tech, Department of Electrical & Computer Engineering	Mar. 2023
Washington University in St. Louis, Department of Computer Science & Engineering	Mar. 2023
Purdue University, School of Electrical & Computer Engineering	Mar. 2023
Purdue University, Department of Computer Science	Mar. 2023
Virginia Tech, Department of Computer Science	Mar. 2023
Indiana University Bloomington, Department of Computer Science	Feb. 2023
University of Colorado Boulder, Department of Computer Science	Feb. 2023
University of Massachusetts Amherst, College of Information and Computer Sciences	Feb. 2023

NURD: Negative-Unlabeled Learning for Online Datacenter Straggler Prediction

Conference presentation at MLSys, Santa Clara, USA	Aug. 2022
--	-----------

Predictable Maintenance Job Planning in Datacenters

Meta Infrastructure Data Science Faculty Workshop at KDD, DC, USA	Aug. 2022
---	-----------

Generalizable and Interpretable Learning for Configuration Extrapolation

Conference presentation at ESEC/FSE, Virtual	Nov. 2021
--	-----------

Dynamical Systems Theory for Causal Inference with Application to Synthetic Controls

Causal Data Science Meeting, Virtual	Nov. 2020
Conference presentation at AISTATS, Virtual	Aug. 2020

Generative and Multi-phase Learning for Computer Systems Optimization

Conference presentation at ISCA, Phoenix, USA	Jun. 2019
---	-----------

Multiresolution Kernel Approximation for Gaussian Process Regression

Conference presentation at NeurIPS, Long Beach, USA

Dec. 2017

Large Scale Kernel Methods for Online AUC Maximization

Conference presentation at ICDM, New Orleans, USA

Nov. 2017

An Adaptive Gradient Method for Online AUC Maximization

Conference presentation at AAAI, Austin, USA

Jan. 2015

Last updated August 22, 2024