

Generative and Multi-phase Learning for Computer Systems Optimization

Yi Ding, Nikita Mishra, Henry Hoffmann



THE UNIVERSITY OF
CHICAGO

Computer Systems Optimization

- Optimizing modern computer systems requires **tradeoffs**:
 - Deliver reliable performance
 - Minimize energy consumption

Computer Systems Optimization


- Optimizing modern computer systems requires **tradeoffs**:
 - Deliver reliable performance
 - Minimize energy consumption
- Resource management via system configuration:
 - Resources have complex, non-linear effects on **performance** and **energy**
 - Resource interactions create **local** optima

Computer Systems Optimization


- Optimizing modern computer systems requires **tradeoffs**:
 - Deliver reliable performance
 - Minimize energy consumption
- Resource management via system configuration:
 - Resources have complex, non-linear effects on **performance** and **energy**
 - Resource interactions create **local** optima
- How to find the optimal **system configuration**?

Example of a Configuration Space \mathcal{C}


$$\mathcal{C} \leftarrow \{\text{Core assignment}\} \times \{\text{Clock speed assignment}\} \times \{\text{Memory controller}\}$$



Clock Speed



Memory controller



Cores

Machine Learning to the Rescue


Machine Learning to the Rescue

- But...
 - **Scarce data**: expensive collection, limited range behavior

Machine Learning to the Rescue

- But...


- **Scarce data**: expensive collection, limited range behavior
- **Asymmetric benefits**: only configs on **optimal frontier** useful



Machine Learning to the Rescue

- But...

- **Scarce data**: expensive collection, limited range behavior → **Generative model**
- **Asymmetric benefits**: only configs on **optimal frontier** useful → **Multi-phase sampling**



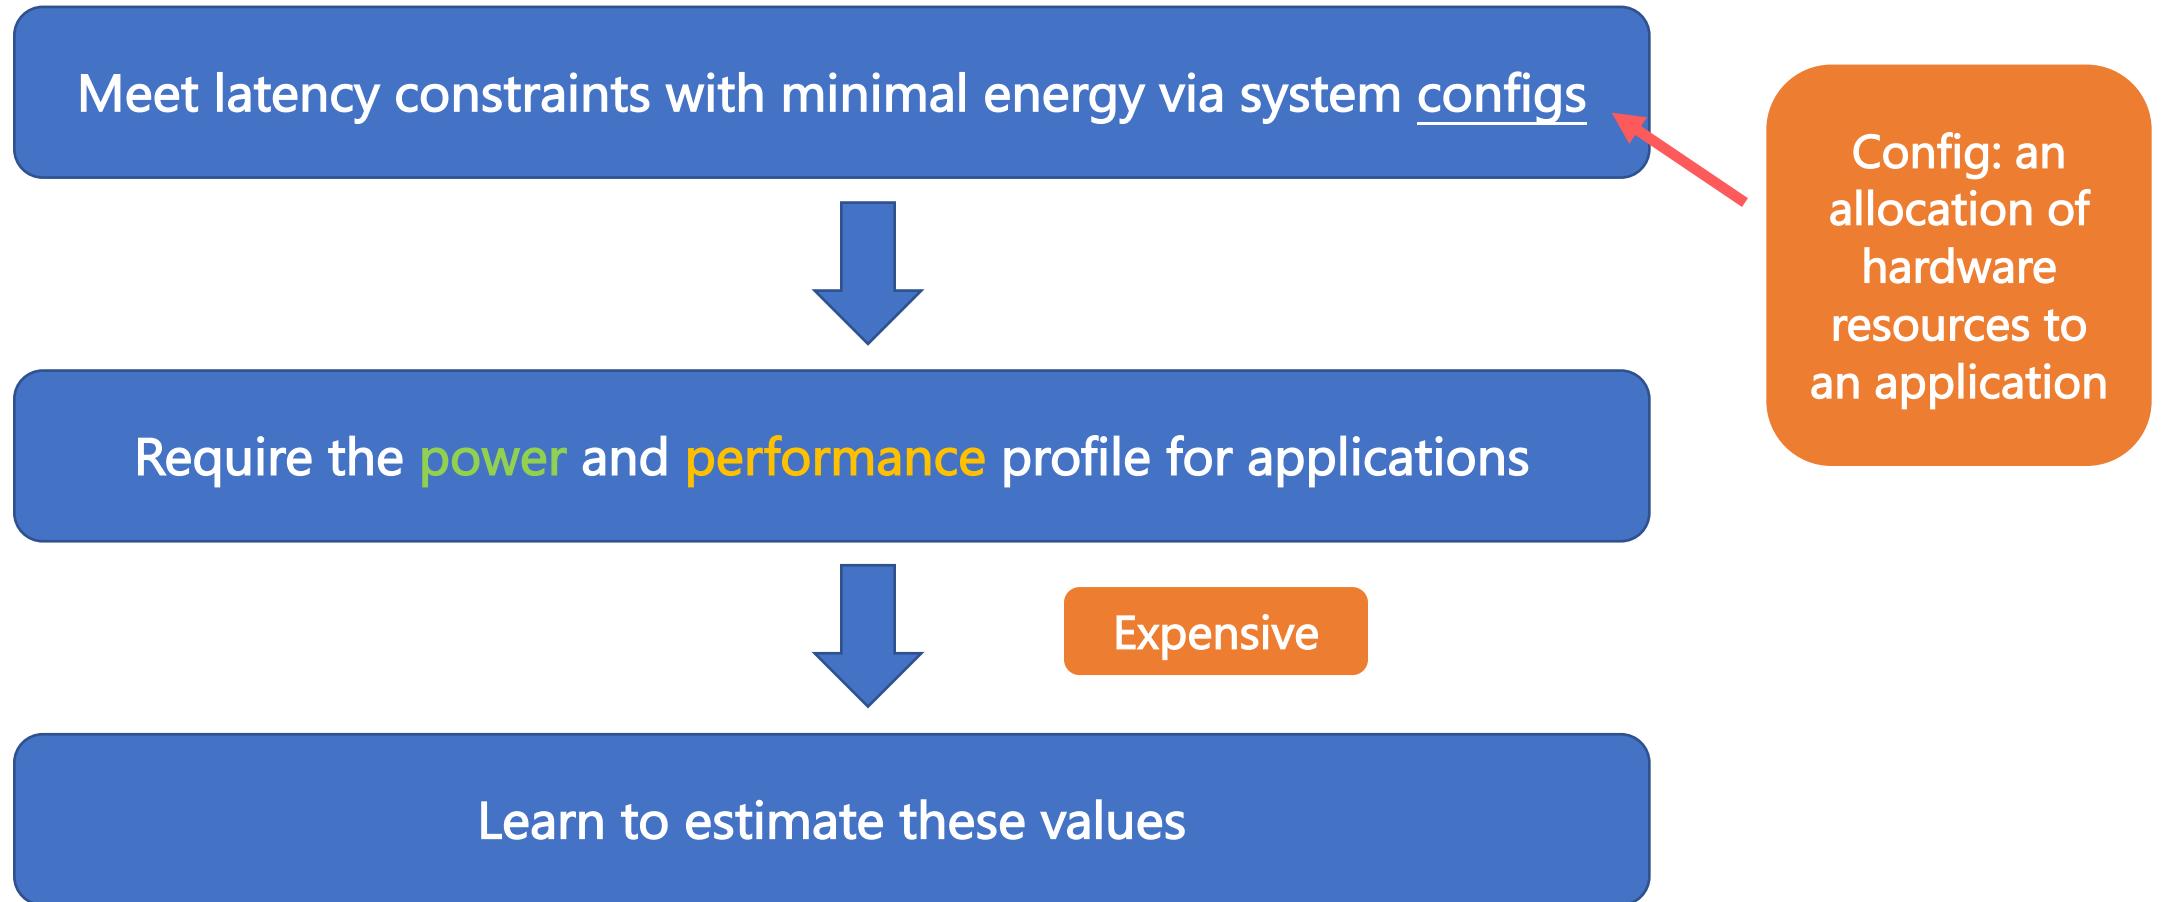
Machine Learning to the Rescue

- But...


- **Scarce data**: expensive collection, limited range behavior → **Generative model**
- **Asymmetric benefits**: only configs on **optimal frontier** useful → **Multi-phase sampling**



Problem Formulation




SRAD on ARM big.LITTLE system



| | Model A | Model B |
|---------------------|--|--|
| Optimal points | Just far enough | True data |
| Non-optimal points | True data | Very far |
| Goodness of fit | 99% | 0 |
| Energy over optimal | 22% ✗ | 0 ✓ |


SRAD on ARM big.LITTLE system



Key Insight:
High accuracy **!=** good system results

| | | |
|---------------------|-----------------|-----------------|
| Optimal points | Just far enough | True data |
| Non-optimal points | True data | Just far enough |
| Goodness of fit | 99% | 0 |
| Energy over optimal | 22% ✗ | 0 ✓ |

Recommender Systems -> Learning by Examples



1. Paragon: QoS-Aware Scheduling for Heterogeneous Datacenters. Christina Delimitrou and Christos Kozyrakis. (ASPLOS 2013)
2. Quasar: Resource-Efficient and QoS-Aware Cluster Management. Christina Delimitrou and Christos Kozyrakis (ASPLOS 2014)

<https://www.le-vpn.com/netflix-10-innovations-changed-world/>
<https://www.thrillist.com/entertainment/nation/the-netflix-prize>

An Analogy

Users

Movies

| | | | | | | | |
|---|---|---|---|---|---|---|---|
| 1 | 3 | ? | 5 | 5 | | | 4 |
| | 5 | 4 | 4 | 2 | | 1 | 2 |
| 2 | | 4 | | 1 | 2 | 3 | |
| 4 | 3 | | 5 | | 2 | 4 | 5 |
| | 4 | 2 | 4 | 2 | 4 | | 1 |
| 2 | | 4 | 1 | | 3 | 2 | 3 |
| 3 | 4 | | 2 | 2 | | 5 | 3 |
| 1 | 3 | 3 | | 2 | 4 | 2 | |



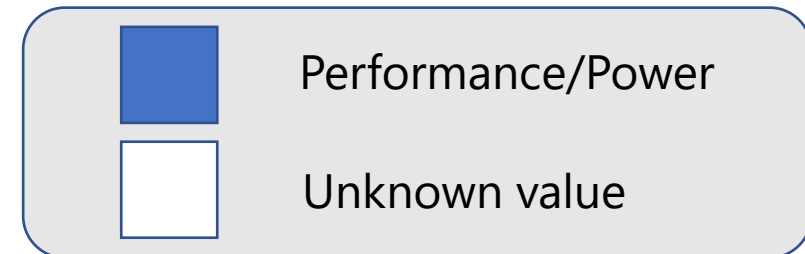
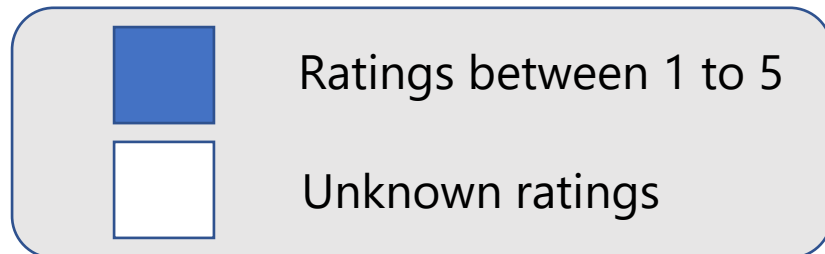
Ratings between 1 to 5



Unknown ratings

An Analogy

| | | Users | | | | | | | |
|--------|---|-------|---|---|---|---|---|---|---|
| Movies | 1 | 1 | 3 | ? | 5 | 5 | | | 4 |
| | 2 | | 5 | 4 | 4 | 2 | | 1 | 2 |
| | 3 | 2 | | 4 | | 1 | 2 | 3 | |
| | 4 | 4 | 3 | | 5 | | 2 | 4 | 5 |
| | 5 | | 4 | 2 | 4 | 2 | 4 | | 1 |
| | 6 | 2 | | 4 | 1 | | 3 | 2 | 3 |
| | 7 | 3 | 4 | | 2 | 2 | | 5 | 3 |
| | 8 | 1 | 3 | 3 | | 2 | 4 | 2 | |

[illegible]

Outline

- Motivation
- **Methods**
- Experimental Results
- Conclusion

Generating Data for Accuracy

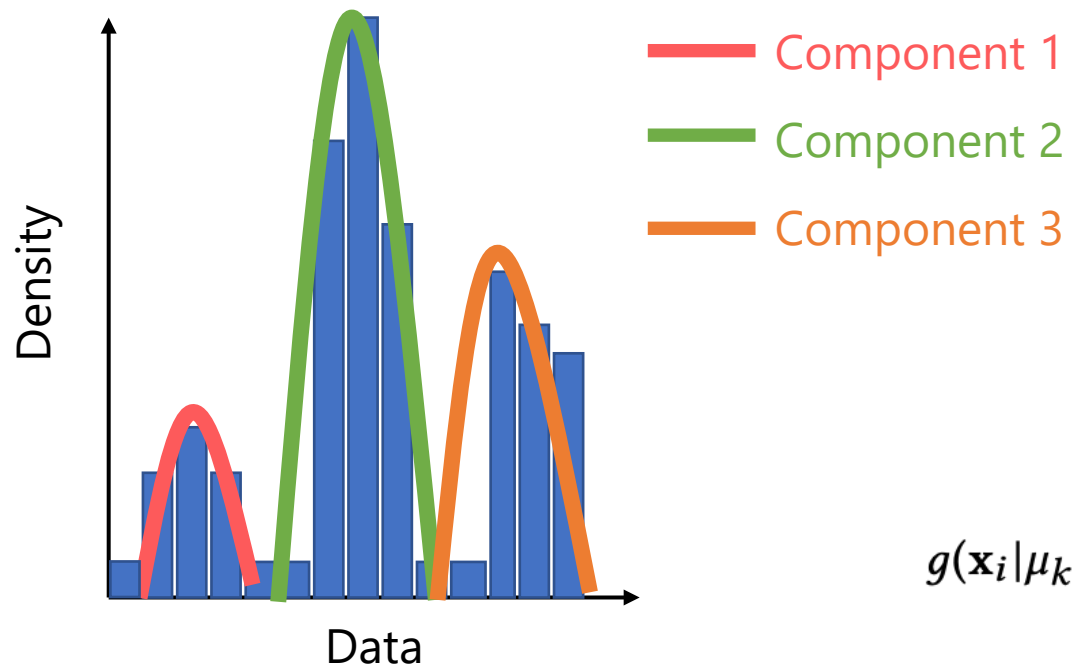
- Goal: **different** enough but still **realistic** to be plausible

Generating Data for Accuracy

- Goal: **different** enough but still **realistic** to be plausible
- How:
 - Random number generator → different but not plausible

Generating Data for Accuracy

- Goal: **different** enough but still **realistic** to be plausible
- How:
 - Random number generator → different but not plausible
 - Gaussian Mixture Model (GMM) → plausible but not different



K : number of components

\mathbf{x}_i : data points, $i=1, \dots, N$

w_k : weight of k -th component

Probability that \mathbf{x}_i belongs to k -th comp:


$$p(\mathbf{x}_i) = \sum_{k=1}^K w_k g(\mathbf{x}_i | \mu_k, \Sigma_k)$$

$$g(\mathbf{x}_i | \mu_k, \Sigma_k) = \frac{1}{(2\pi)^{\frac{D}{2}} |\Sigma_k|^{\frac{1}{2}}} \exp \left(-\frac{1}{2} (\mathbf{x}_i - \mu_k)^\top \Sigma_k^{-1} (\mathbf{x}_i - \mu_k) \right)$$

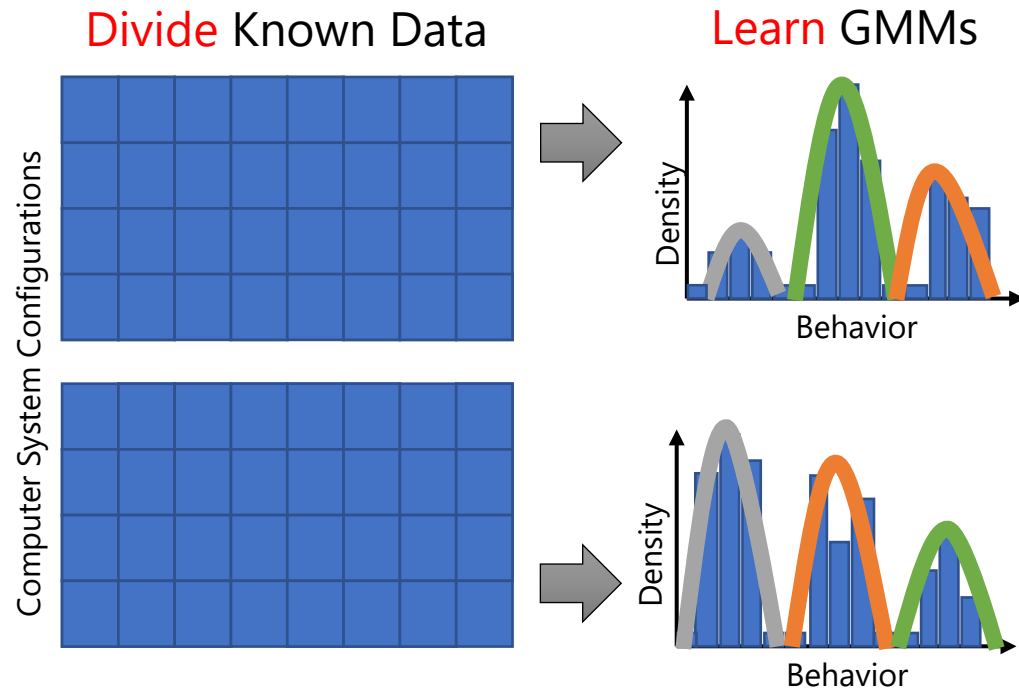
Generating Data with a GMM

Divide Known Data

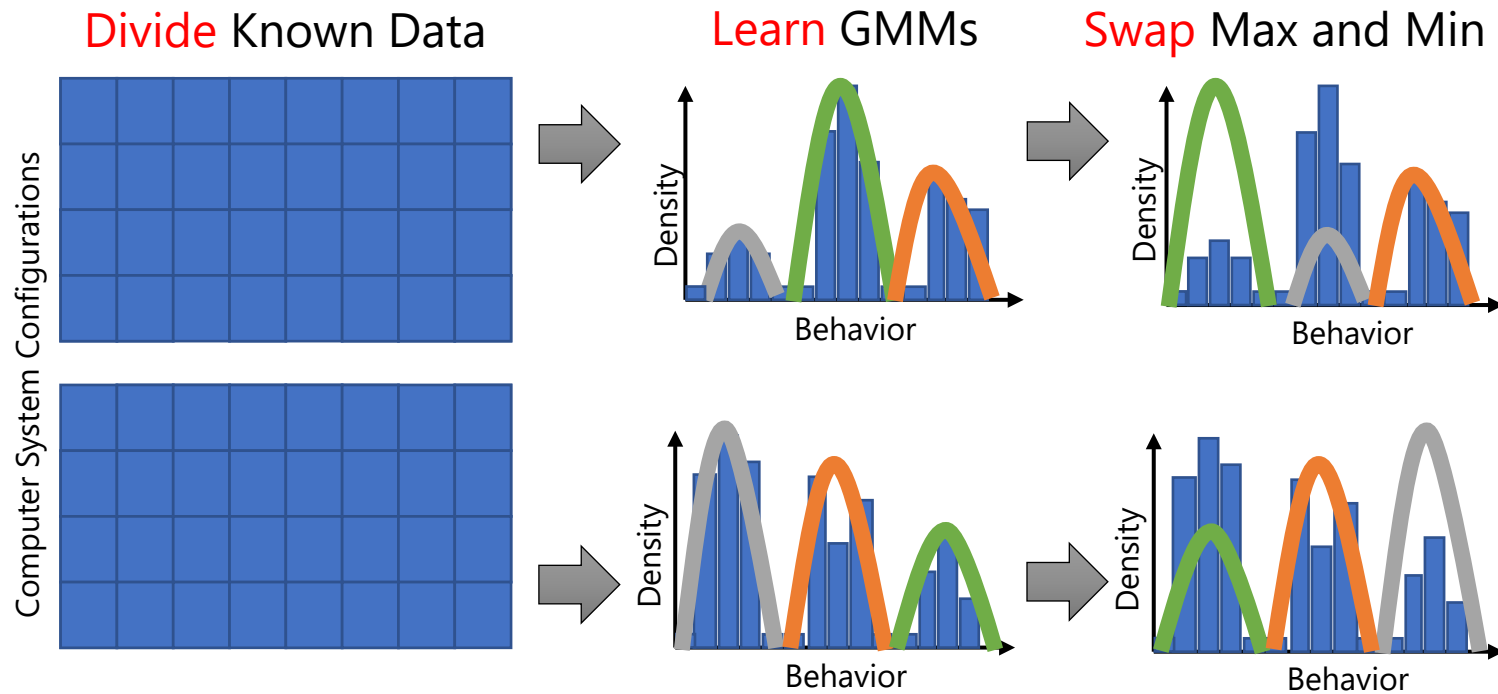
Computer System Configurations



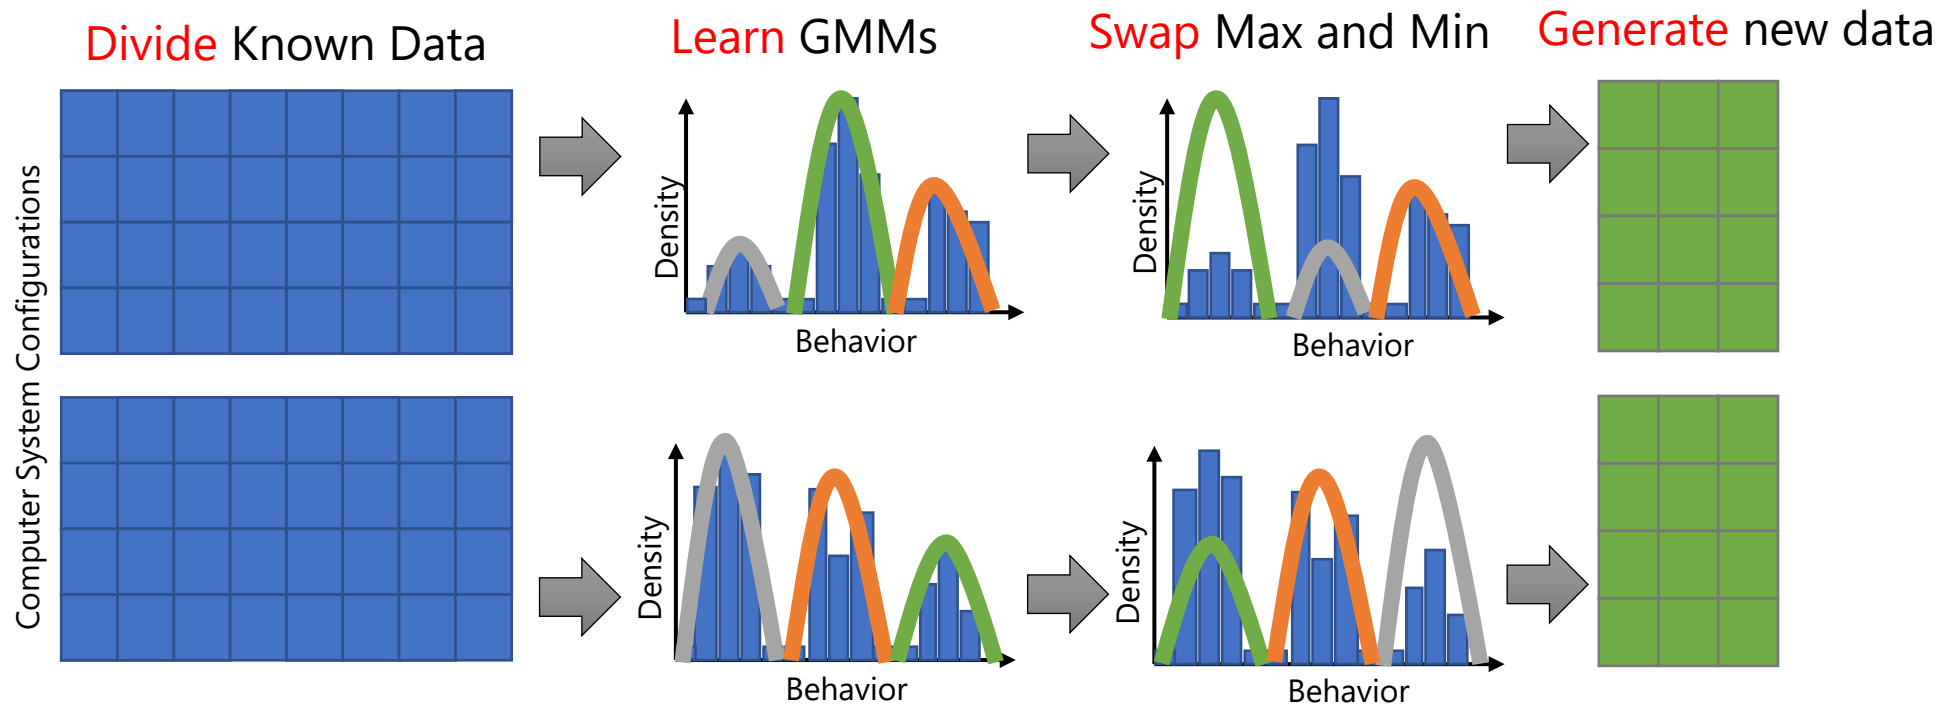
Generating Data with a GMM



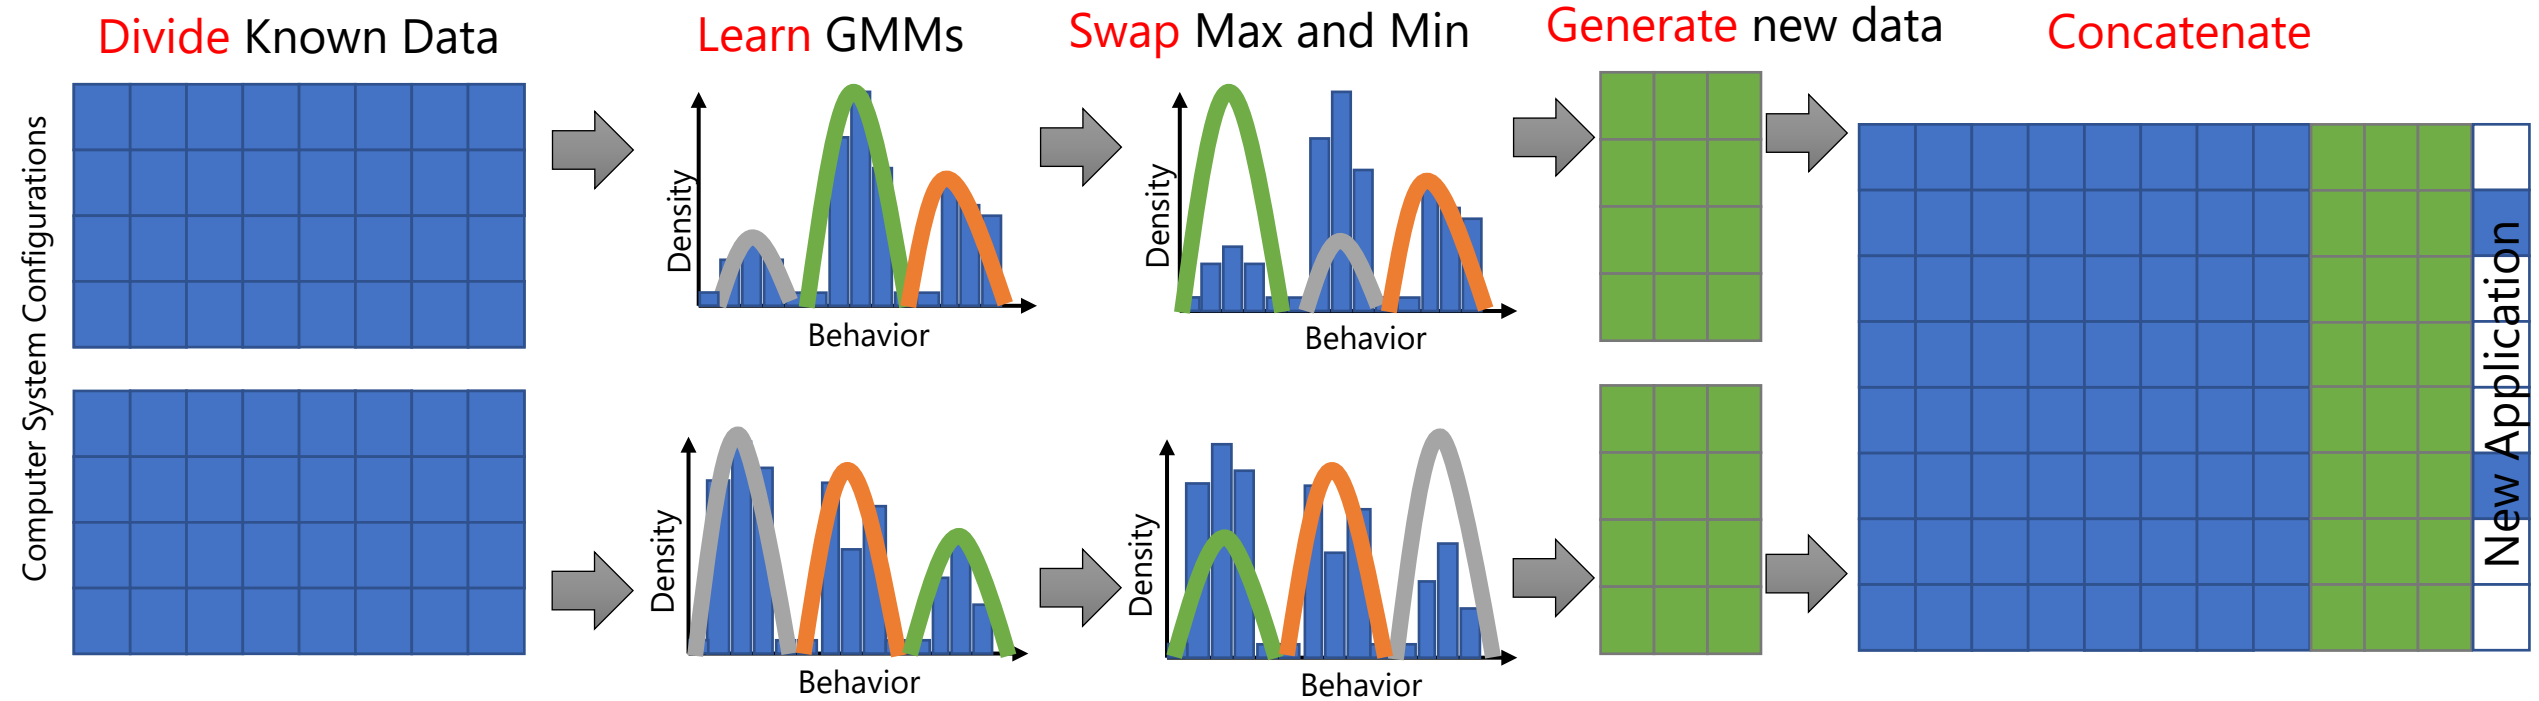
Generating Data with a GMM



Generating Data with a GMM



Generating Data with a GMM



Multi-phase Sampling

Input: Configuration-Application data matrix, Sampling budget N

Matrix Completion with Sample Size $N/2$

Known Applications


[illegible]

Multi-phase Sampling

Input: Configuration-Application data matrix, Sampling budget N

Matrix Completion
with Sample Size $N/2$
Known Applications

Estimated
Behavior for New
Application




Multi-phase Sampling

Input: Configuration-Application data matrix, Sampling budget N

Matrix Completion
with Sample Size $N/2$
Known Applications

Estimated Behavior for New Application

Select $N/2$
Best Configs



$$\text{efficiency} = \frac{\text{estimated performance}}{\text{estimated power}}$$

Multi-phase Sampling

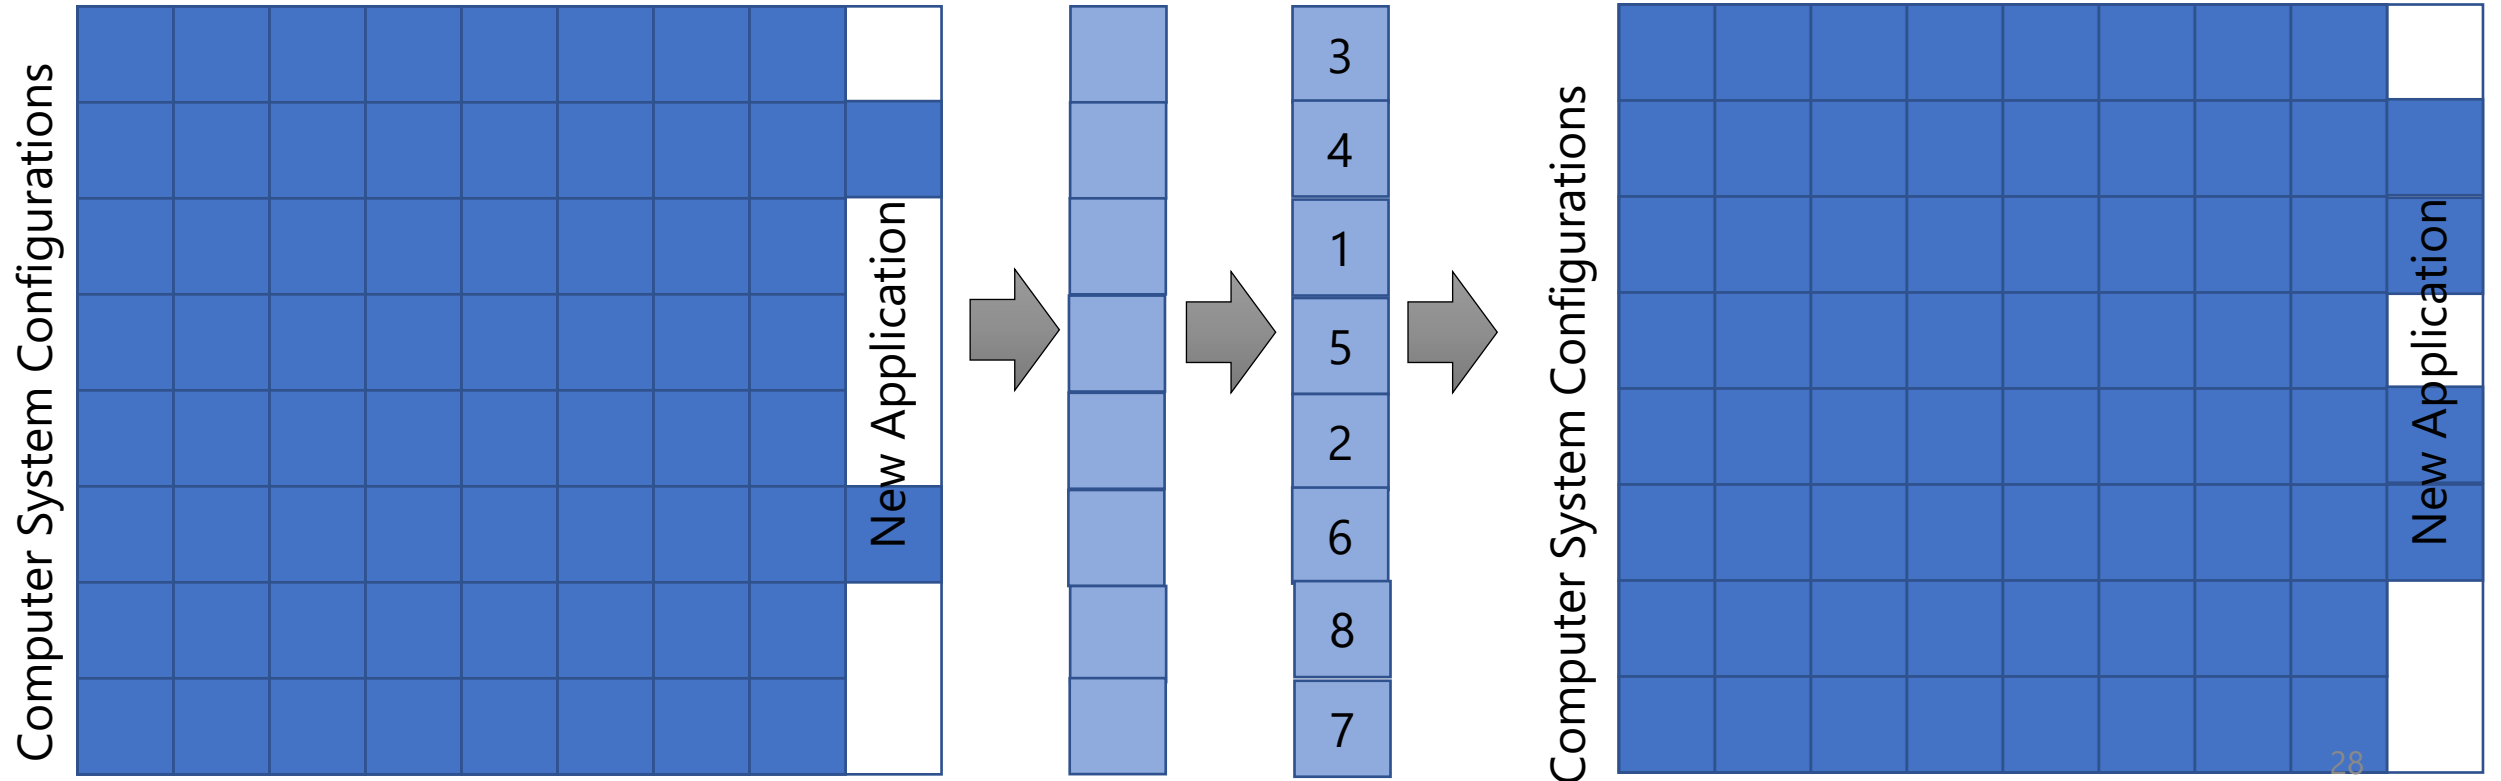
Input: Configuration-Application data matrix, Sampling budget N

Matrix Completion
with Sample Size $N/2$
Known Applications

Estimated
Behavior for New
Application

Select $N/2$
Best Configs

Matrix Completion with $N/2$ original
samples and $N/2$ estimated best configs
Known Applications




Outline

- Motivation
- Methods
- Experimental Results
- Conclusion

Experimental Setup

| | Mobile | Server |
|------------------|----------------|--------------------|
| System | Ubuntu 14.04 | Linux 3.2.0 system |
| Architecture | ARM big.LITTLE | Intel Xeon E5-2690 |
| # Applications | 21 | 22 |
| # Configurations | 128 | 1024 |



Learning Models and Frameworks

| Learning Models | Category |
|-----------------|----------|
| MCGD | MC |
| MCMF | MC |
| Nuclear | MC |
| WNNM | MC |
| HBM | Bayesian |

First comprehensive study of matrix completion (MC) algorithms for systems optimization task


Learning Models and Frameworks

| Learning Models | Category |
|-----------------|----------|
| MCGD | MC |
| MCMF | MC |
| Nuclear | MC |
| WNNM | MC |
| HBM | Bayesian |


First comprehensive study of matrix completion (MC) algorithms for systems optimization task

| Frameworks | Definitions |
|------------|----------------------|
| Vanilla | Basic learners |
| GM | Generative model |
| MP | Multi-phase sampling |
| MP-GM | Combine GM and MP |

Improve Prediction Accuracy w/ GM



Mobile



Server

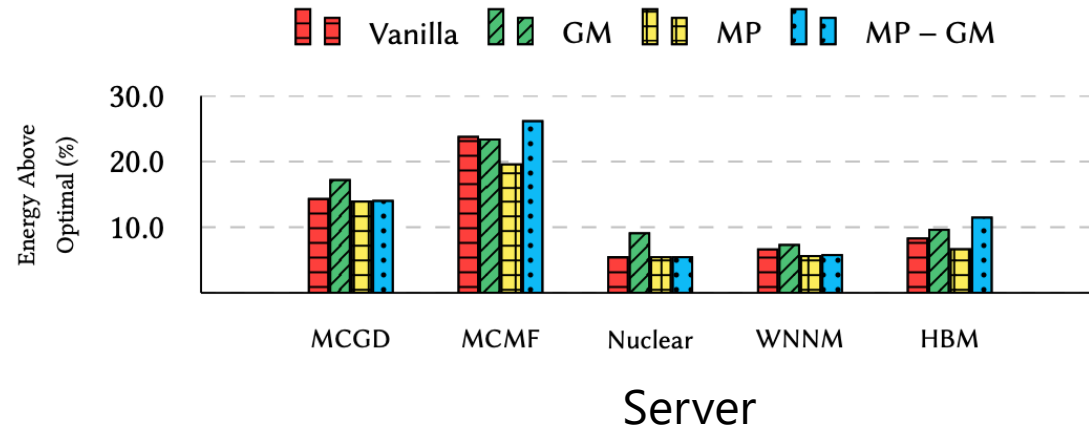
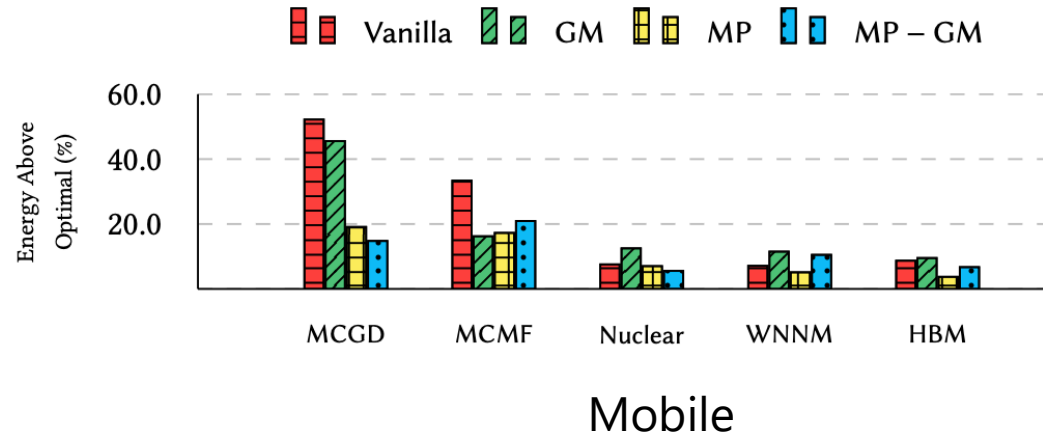


High
is
Better

Average percentage points of accuracy improvement

| | | GM | MP | MP – GM |
|----------------|-------------|------|------|---------|
| Mobile | Performance | 1.8 | 1.4 | 2.3 |
| | Power | 4.3 | 0.6 | 3.4 |
| Server | Performance | 9.0 | -0.2 | -0.3 |
| | Power | 20.5 | -0.4 | 0.1 |
| Average | | 8.9 | 0.4 | 1.4 |

Improve Energy Savings w/ MP



Lower
is
Better


Average energy improvement

| | GM | MP | MP – GM |
|--------|------|-----|---------|
| Mobile | –14% | 41% | 22% |
| Server | –22% | 11% | –6.5% |


Conclusion



Clock Speed






Memory controller







Cores

Conclusion




Conclusion




Conclusion




Clock Speed



Memory controller



Cores



We advocate:

- De-emphasizing prediction accuracy
- Incorporating system structure into learner

Yi Ding, Nikita Mishra, and Henry Hoffmann. 2019. Generative and Multiphase Learning for Computer Systems Optimization. In The 46th Annual International Symposium on Computer Architecture (ISCA '19)